

# Dynamic Pricing in Transportation Network Systems as a Markov Decision Process

Matt Battifarano

12-735 Urban Systems Modeling

7 May 2018

# Outline

Introduction

Some Simplifying Assumptions

The Model

An Example Transition

Finding the Optimal Policy

Results

Future Work

Questions?

# Introduction

- ▶ TNCs like Uber and Lyft use **surge pricing** to balance supply and demand.
- ▶ TNCs want to choose the surge multiplier which **maximizes revenue** given the supply and demand at a location.
- ▶ A **Markov Decision Process** can be used to determine the **optimal surge multiplier policy**.

# Assumptions

- ▶ Single location
- ▶ The **state** at time  $t$  is given by the **number of users** and the **number of drivers**:  $s_t = (u_t, d_t)$
- ▶ Users *always* accept a ride if there is no surge multiplier.
- ▶ Drivers *always* accept users.
- ▶ The **action** at time  $t$  is the **surge multiplier**.
- ▶ When the surge multiplier is **high**:
  - ▶ **Users** will be more likely to **wait**.
  - ▶ Some may decide to **leave**.
  - ▶ **Drivers** will be **attracted** to the location.
- ▶ **Users** and **drivers** arrive stochastically at the location with rates  $\lambda_u$  and  $\lambda_d(m)$ .
- ▶ User and driver choices are i.i.d.


# The Model

Variables:

- ▶  $u_t$ : number of users
- ▶  $d_t$ : number of drivers
- ▶  $m_t$ : surge multiplier
- ▶  $r_t$ : revenue

Transitions:


- ▶  $u_{t+1} = u_t^{(\text{remain})} + \tilde{u}_{t+1}$
- ▶  $d_{t+1} = d_t - d_t^{(\text{match})} + \tilde{d}_{t+1}$



# The Model

User Transition:

$$u_{t+1} = u_t^{(\text{remain})} + \tilde{u}_{t+1}$$



- ▶  $u_t^{(\text{accept})} \sim \text{Binom}(p(\text{accept} \mid m_t), u_t)$
- ▶  $u_t^{(\text{match})} = \min(u_t^{(\text{accept})}, d_t)$
- ▶  $u_t^{(\text{remain})} \sim \text{Binom}(p(\text{stay}), u_t - u_t^{(\text{match})})$
- ▶  $\tilde{u}_{t+1} \sim \text{Poisson}(\lambda_u)$


# The Model

Driver Transition:

$$d_{t+1} = d_t - d_t^{(\text{match})} + \tilde{d}_{t+1}$$

►  $d_t^{(\text{match})} = u_t^{(\text{match})}$

►  $\tilde{d}_{t+1} \sim \text{Poisson}(\lambda_d(m_t))$




# The Model

Revenue:

$$r_t = R(s_t, m_t) = \text{fares} - \text{opportunity cost}$$

- ▶  $\text{fares} = d_t^{(\text{match})} \cdot m_t \cdot \text{base fare}$
- ▶  $\text{opportunity cost} = (d_t - d_t^{(\text{match})}) \cdot \text{penalty}$
- ▶ Opportunity cost penalizes empty vehicles.

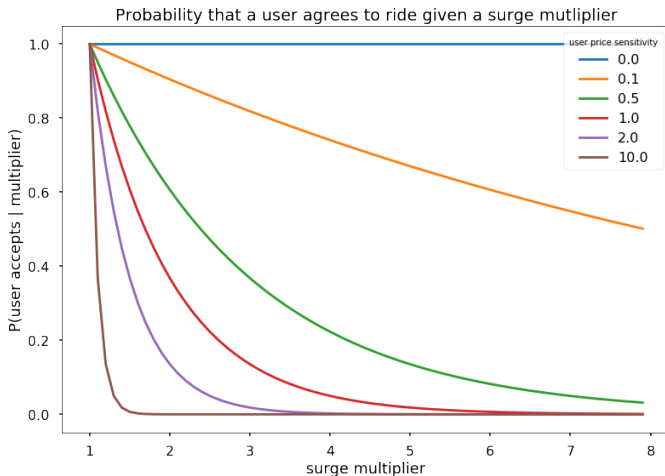




## An Example Transition

What is the probability that a user decides to take a ride *after* seeing the multiplier?

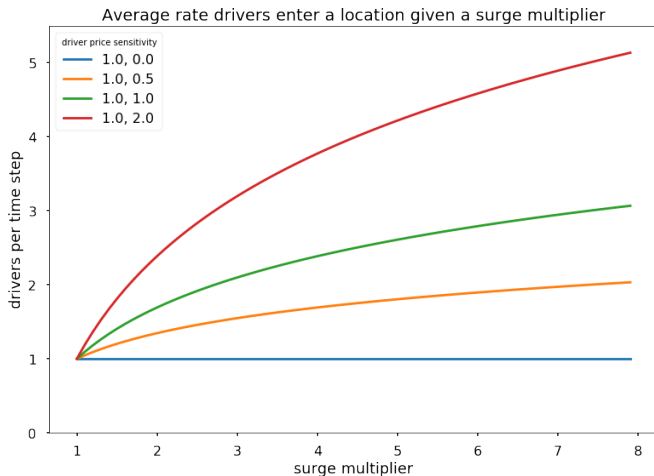
$$p(\text{accept} \mid m) = \exp(-\alpha(m - 1))$$



## An Example Transition


How many drivers do we expect to arrive to the location *after* seeing the multiplier?

$$\lambda_d(m) = \beta_1 \log(m) + \beta_0$$




# An Example Transition

What if drivers and riders **are balanced**?




# An Example Transition

What if drivers **are abundant**?



# An Example Transition

What if drivers **are scarce**?



# Finding the Optimal Policy

## Value Iteration

- ▶ Initialize the value vector  $V_0$
- ▶ For  $k = 1, \dots$ 
  - ▶ For each state  $s$ , compute the optimal policy:

$$\pi_k(s) = \arg \max_m \mathbb{E}_{s' \sim p(s_{t+1}=s' | s_t=s, m_t=s)} [R(s_t, m_t) - \gamma V_k(s')]$$

- ▶ For each state  $s$ , update  $V$  using  $\pi_k$

$$V_{k+1}(s) = \mathbb{E}_{s' \sim p(s_{t+1}=s' | s_t=s, m_t=\pi_k(s))} [R(s_t, \pi_k(s)) - \gamma V_k(s')]$$


- ▶ Until  $V_{k+1} \approx V_k$ .
  - ▶ Return  $\pi = \pi_{k+1}$

# Finding the Optimal Policy

$$\mathbb{E}_{s' \sim p(s_{t+1}=s' | s_t=s, m_t=s)}[R(s_t, m_t) - \gamma V_k(s')]$$

- ▶ Problem:  $p(s_{t+1} | s_t, m_t)$  is not known in closed form.
- ▶ Fortunately, it is **easy to sample from**.
- ▶ Idea: **approximate** the expectation with **Monte Carlo sampling**.


# Results






## Results

What does the optimal policy look like when  $\alpha = 2$ ?



## Results

What about when users **care less** about surges? ( $\alpha = 0.5$ )



## Future Work

- ▶ Make larger state spaces traceable: Approximate  $V_k$  and  $\pi_k$  using sampling and a **Gaussian Process**.
- ▶ Can we solve the **inverse problem**: given a sequence of surge multipliers can we infer the parameters  $(\alpha, \beta_1, \beta_0)$ ?

# Questions?

Thank you!

Code and Jupyter Notebook are available on GitHub:

<https://github.com/mbattifarano/surge-multiplier-mdp>